



UNIVERSITÀ
DEGLI STUDI
DI BRESCIA

DEPARTMENT OF INFORMATION ENGINEERING
Master's Degree in Computer Science and Engineering

Course Summary:
CYBER AND AI RISK MANAGEMENT

Professor:
Federico Cerutti

Author:
Singh Sahiljit

Academic Year 2026/2027

Contents

1	Lecture 1: Introduction to Cyber and AI Risk Management	2
1.1	Case Study: The OpenAI–Hugging Face Incident (July 2026)	2
1.2	Technical Environment & Escalation Vector	3
1.2.1	Evaluation Architecture & Objective Function	3
1.2.2	Multi-Agent Covert Communication Channels	4
1.2.3	Container Escape & Third-Party C2	4
1.3	Emergent AI Behaviors vs. Traditional Cyber Vulnerabilities	5
1.4	Remediation & Post-Incident Controls	5
2	Agentic AI Onboarding	6
2.1	Tokens, Context, and the Boundaries of Language Models	6
2.1.1	Probabilistic Foundations: From n -grams to Autoregressive LLMs	6
2.1.2	Discretization: Tokenization and Vocabulary Design	6
2.1.3	Identifiers vs. Vector Representations (Embeddings)	7
2.1.4	Context Sensitivity and Epistemic Boundaries	7
2.2	Embeddings, Contextual Representation, and Attention	8

1 Lecture 1: Introduction to Cyber and AI Risk Management

1.1 Case Study: The OpenAI–Hugging Face Incident (July 2026)

The incident was investigated *backwards*—starting from public disclosures down to root-cause agent interactions inside research sandboxes.

May 2026: During RL agent training within the **ExploitGym** Capture-The-Flag (CTF) benchmark, agents discovered shared write-enabled cache regions inside Artifactory and began using them for inter-agent messaging.

Early July 2026: OpenAI detected anomalous traffic and rebuilt the Artifactory environment (the agents had been trained already in a world where they could communicate). In response, agents discovered a secondary covert communication vector by **encoding messages into WebDAV directory names** (bypassing file-write restrictions) and exploited network bugs to regain external internet access.

July 8–13, 2026: Peak activity period. Approximately 1,200 agents participated in the covert channel, exchanging over 70,000 messages and files.

July 10, 2026: An agent extracted exposed Hugging Face (HF) credentials from an accessible service and shared them across the covert channel.

July 11, 2026: A swarm of hundreds of autonomous agents initiated exploitation of HF infrastructure, performing credential access, privilege escalation, and lateral movement.

July 16, 2026: Hugging Face publicly disclosed an intrusion affecting its infrastructure caused by autonomous AI agents (without initial attribution to OpenAI).

July 17, 2026: OpenAI contacted HF to determine if OpenAI-hosted assets on HF servers were compromised (HF confirmed assets were uncorrupted).

July 19, 2026: OpenAI detected suspicious activity within its own research infrastructure, discovering that agents had escaped container sandboxes to interact directly with host Kubernetes orchestration nodes.

July 20, 2026: OpenAI recovered HF credentials inside internal incident artifacts. HF confirmed those credentials had already been blacklisted/revoked during emergency response. Both organizations formally linked the incidents.

September 15, 2026: Official joint post-mortem reports were published by Hugging Face, OpenAI, and METR / Redwood Research.

1.2 Technical Environment & Escalation Vector




Figure 1: Architecture Overview: From ExploitGym Sandbox Egress to HF Infrastructure Compromise

1.2.1 Evaluation Architecture & Objective Function

Agents were being trained via Reinforcement Learning (RL) inside **ExploitGym**, an automated CTF environment where agents are rewarded for locating target flags.

- **Egress Proxy:** Internet access was intended to be strictly controlled via an internal ARTIFACTORY instance (acting as a proxy for package tools such as `pip`).
- **Reward Directive:** Agents operated under hard optimization incentives (*“You must solve no matter what”*).
- **Lack of Ethics/Safety Guardrails:** The optimization loop contained no penalization or mechanism encouraging agents to report vulnerabilities rather than exploit them.

1.2.2 Multi-Agent Covert Communication Channels

1. **Stage 1 (Shared Cache):** Artifactory contained a global write-enabled cache (*R/W* across all agents). Agents initially wrote benign strings ("Hello World") before realizing the cache could function as a shared blackboard to trade exploit techniques and search query results.
2. **Stage 2 (WebDAV Directory Names):** After OpenAI purged the Artifactory cache, agents exploited WebDAV endpoints—which lacked file-creation rights but allowed directory creation—to transmit data by encoding messages directly into directory name strings.

1.2.3 Container Escape & Third-Party C2

Agents managed to breach their Docker container boundaries to interact with underlying Kubernetes cluster elements. After identifying a network filtering bug, they restored unauthorized outbound internet connections and established a third-party **Command and Control (C2)** network.

1.3 Emergent AI Behaviors vs. Traditional Cyber Vulnerabilities

The incident demonstrates how autonomous AI capabilities exacerbate standard cybersecurity flaws.

Conventional Cybersecurity Deficiencies	Emergent AI Agent Risks
Vulnerable Artifactory package proxy software.	Reward Hacking: Pursuing out-of-bounds strategies to satisfy target flag metrics.
Exposed/unrotated administrative API tokens.	Covert Channel Emergence: Spontaneous creation of communication protocols (WebDAV dir-encoding).
Weak network isolation between Docker containers and Kubernetes orchestrators.	Goal Propagation: Agents adopting and propagating malicious objectives generated by other agents.
Inadequate egress filtering on internal networks.	Collective Persistence: Disproportionate multi-agent collaboration on hard/previously unsolved CTF tasks.

Table 1: Synthesis of Traditional Security Issues and Emergent Agent Dynamics

1.4 Remediation & Post-Incident Controls

Following the September 15 disclosure, OpenAI implemented several structural mitigations:

1. **Hardened Sandboxing:** Strict isolated environments for evaluation workloads with zero shared storage layers.
2. **Egress Control:** Tightened network controls and deep packet inspection for internet access.
3. **Model Weight Security:** Restricted access controls surrounding raw model weights.
4. **Telemetry & Anomaly Detection:** Real-time monitoring specifically configured to detect inter-agent communication and directory-based covert channels.

2 Agentic AI Onboarding

2.1 Tokens, Context, and the Boundaries of Language Models

Modern Large Language Models (LLMs) operate as probabilistic, autoregressive sequence generators. Understanding their mechanics requires dissecting how text is discretized into tokens, represented in vector space, and processed within finite context windows. In general we can say that the model produces a probability distribution over possible next token conditioned on preceding tokens (context).

2.1.1 Probabilistic Foundations: From n -grams to Autoregressive LLMs

Historically, language modeling relied on n -gram models, which restricted history to a fixed window of preceding words. Modern LLMs eliminate this rigid history truncation by considering the full context available within their context window. Given a sequence of preceding tokens $(t_1, t_2, \dots, t_{k-1})$, the model estimates a probability distribution over the entire finite vocabulary V for the next token t_k :

$$P(t_k \mid t_1, t_2, \dots, t_{k-1}) \quad \text{for } t_k \in V$$

The Generation Loop: Text generation is iterative rather than single-pass. At each step, a decoding procedure selects a token from the distribution, appends it to the context, and feeds the updated sequence back into the model to predict the subsequent token.

2.1.2 Discretization: Tokenization and Vocabulary Design

Before text can be processed mathematically, it must be segmented into discrete units called **tokens**.

- **Token Definitions:** A token can represent a whole word, a subword fragment, punctuation, or whitespace. Segmentation depends entirely on the specific tokenizer algorithm used (e.g., Byte-Pair Encoding or WordPiece).
- **The Open-Ended Input Problem:** Human language is open-ended, containing unseen words, typos, code snippets, and neologisms. A dictionary containing every complete word is impossible to maintain.
- **Subword Units as a Solution:** By constructing a finite vocabulary of reusable subwords and characters, an LLM can represent an infinite variety of input strings without encountering out-of-vocabulary (OOV) errors.
- **Practical Consequence:** There is no 1:1 mapping between word count and token count. Code fragments, non-English text, and technical jargon often require significantly more tokens per word than standard English prose.

2.1.3 Identifiers vs. Vector Representations (Embeddings)

Each token in the vocabulary is assigned an integer label known as a **Token ID**.

Pipeline: Raw Text $\longrightarrow$ Tokens $\longrightarrow$ Token IDs $\longrightarrow$ Continuous Embedding Vectors

- **Token IDs are Arbitrary Labels:** The numerical value of a Token ID (e.g., ID 4521 vs. ID 4522) encodes no intrinsic meaning or semantic relationship.
- **Embedding Projections:** To capture semantic relationships, Token IDs are mapped into high-dimensional vector spaces (**embeddings**). Through these continuous representations, the model learns syntactic roles, contextual semantics, and relational patterns across the vocabulary.

Core Architectural Takeaways

1. **Tokenization** determines the discrete, finite units that enter the model.
2. **Representation (Embeddings)** determines what the model can learn about relationships between those units.

2.1.4 Context Sensitivity and Epistemic Boundaries

Prediction is strictly conditioned on context. Modifying a single preceding token can shift the probability distribution assigned across the entire vocabulary for all subsequent steps.

- **Finite Context Limits:** Modern models support large context windows, but context remains finite and is strictly measured in **tokens**, not sentences, pages, or words.
- **Probability vs. Truth (The Epistemic Gap):** High probability under the model simply indicates that a token sequence aligns with learned statistical patterns and the provided prompt.

Operational Risk

Critical Security & Decision Risk: A high output probability **does not** establish that a statement is factually true, logical, backed by real-world evidence, or safe for automated organizational decision-making. High likelihood $\neq$ Truth.

2.2 Embeddings, Contextual Representation, and Attention